

Assignment 1:

Scale Invariant Feature Detection and Image Filtering

Computer Vision
NTU, Spring 2025

Announced: 2025/03/07 (Fri.)

Due: 2025/03/27 (Thu.) 23:59

Outline

Part 1: Scale Invariant Feature Detection

- Implement Difference of Gaussian (DoG)

Part 2: Image Filtering

- Advanced color-to-gray conversion
- Implement bilateral filter

Rules

- Environment
- Report
- Submission
- Grading
- Late Policy

Outline

Part 1: Scale Invariant Feature Detection

- Implement Difference of Gaussian (DoG)


Part 2: Image Filtering

- Advanced color-to-gray conversion
- Implement bilateral filter



Rules

- Environment
- Report
- Submission
- Grading
- Late Policy


Laplacian of Gaussian



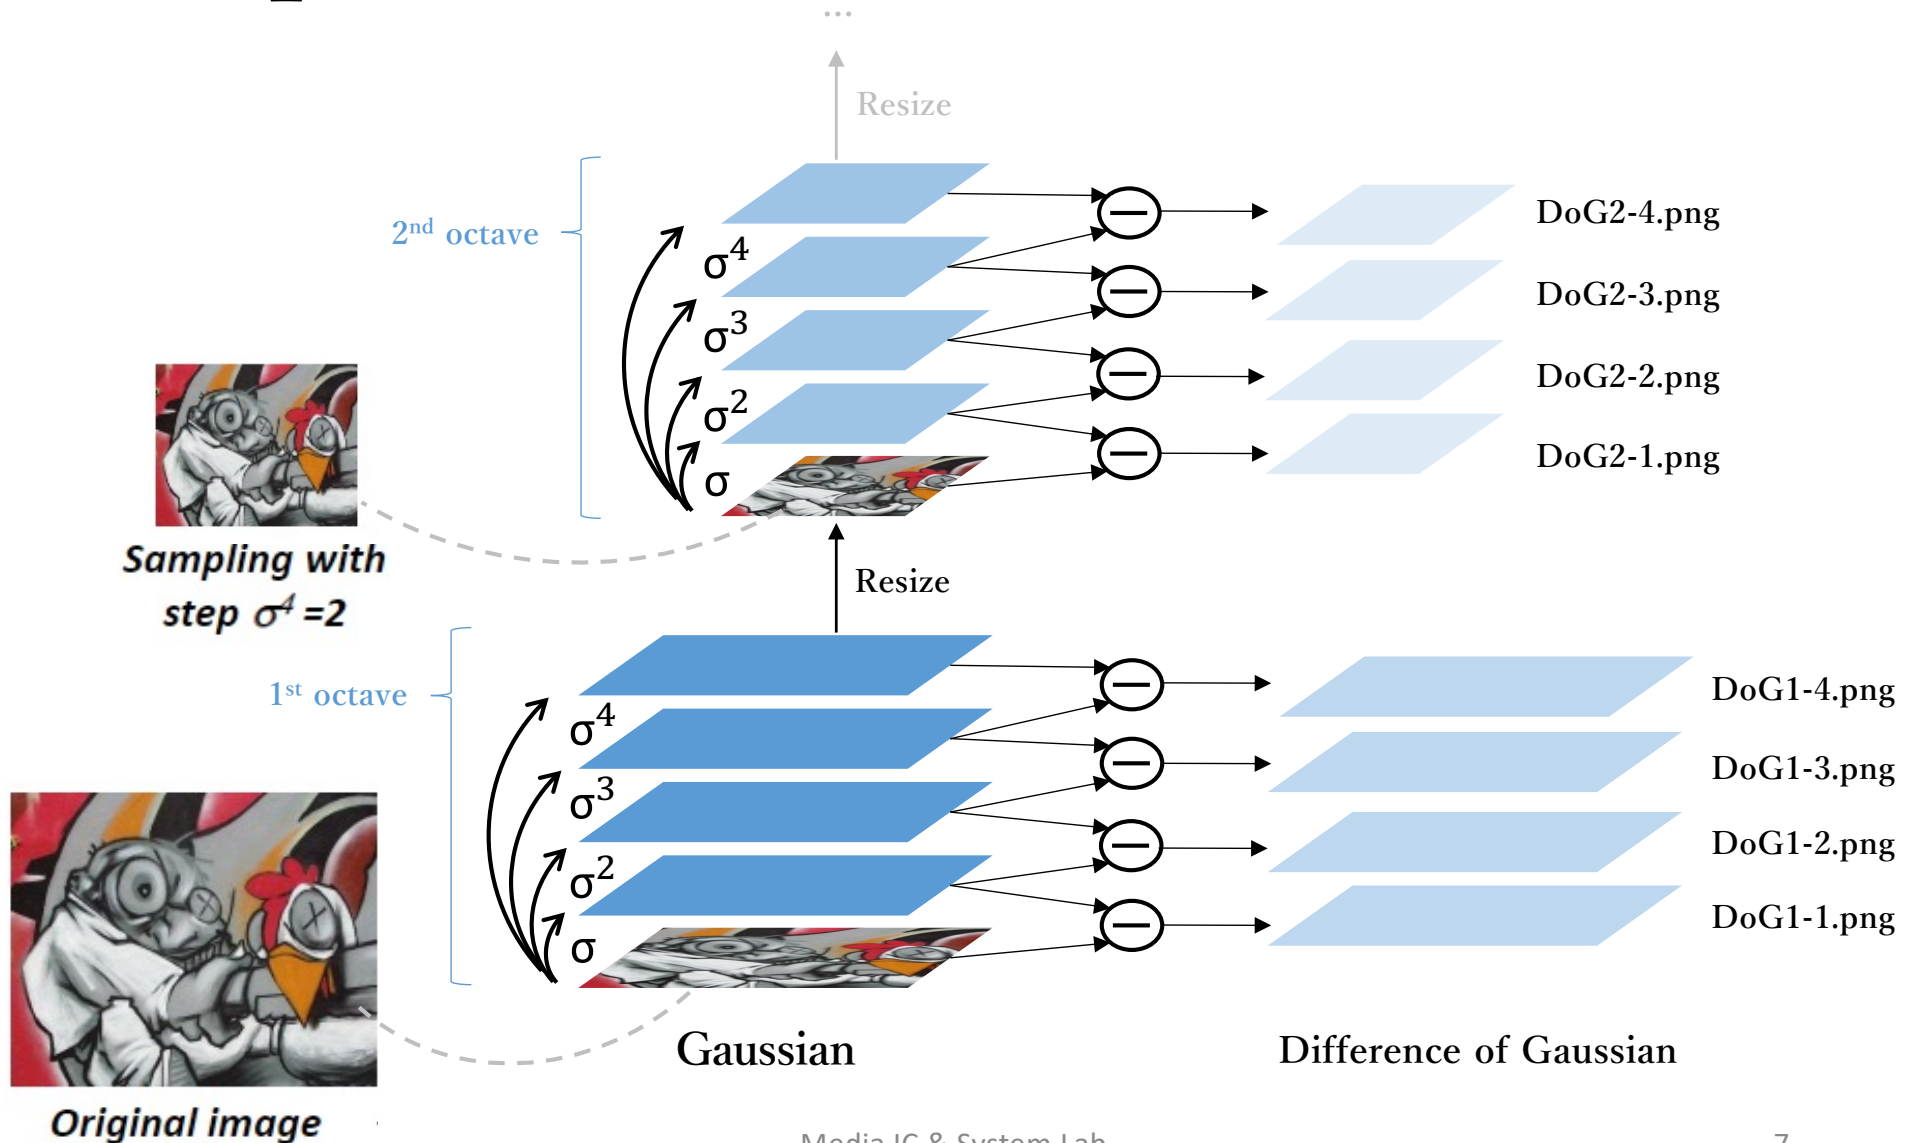
Difference of Gaussian Filter



Gaussian Pyramid



Implementation




Implementation


- Find local extremum (maximum and minimum)
- Padding is not needed



Gaussian Blur
+
Difference



Scale



List of (x, y)
or
(2x, 2y) if resize to 1/2

Implementation

In DoG.py

```
# Step 1: Filter images with different sigma values (5 images per octave, 2 octave in total)
# - Function: cv2.GaussianBlur (kernel = (0, 0), sigma = self.sigma**__)
```

- Apply gaussian blur with corresponding sigma value on the input.
- The base image of the 1st octave is the input image.
- Down sample the last blurred image in the 1st octave as the base image of the 2nd octave.

```
# Step 2: Subtract 2 neighbor images to get DoG images (4 images per octave, 2 octave in total)
# - Function: cv2.subtract(second_image, first_image)
```

- You should subtract the second image (more blurred one) to the first image (less blurred one) to get DoG.

```
# Step 3: Thresholding the value and Find local extremum (local maximum and local minimum)
#         Keep local extremum as a keypoint
```

- Find the local extremum and threshold the pixel value.

```
# Step 4: Delete duplicate keypoints
# - Function: np.unique
```

- Delete duplicate keypoints from 2 octaves.

Assignment Description

- `part1/eval.py`
 - TA will run this code to evaluation your result.
 - **DO NOT Modify!**
- `part1/main.py`
 - Execute DoG, visualize results for report, ... etc.
- `part1/DoG.py`
 - Follow the instructions and implement Difference of Gaussian.
 - The output format should be a numpy array with shape $(x, 2)$
 - x is the number of feature points

Assignment Description

- Recommended steps

1. Implement Difference of Gaussian in DoG.py
2. Use eval.py to evaluate your DoG.py

```
$ python3 eval.py --threshold 3.0 --image_path ./testdata/1.png --gt_path ./testdata/1_gt.npy
```

- Your Result needs to match Ground truth

```
[Info] All keypoints match.
```

3. Finish remaining code in main.py if needed
 - e.g. visualize DoG, plot feature points

Outline

Part 1: Scale Invariant Feature Detection

- Implement Difference of Gaussian (DoG)

Part 2: Image Filtering

- Advanced color-to-gray conversion
- Implement bilateral filter

Rules

- Environment
- Report
- Submission
- Grading
- Late Policy





Color Conversion

- RGB2YUV
 - Read <https://en.wikipedia.org/wiki/YUV> for more details

$$\begin{bmatrix} Y' \\ U \\ V \end{bmatrix} = \begin{bmatrix} 0.299 & 0.587 & 0.114 \\ -0.14713 & -0.28886 & 0.436 \\ 0.615 & -0.51499 & -0.10001 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix},$$
$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = \begin{bmatrix} 1 & 0 & 1.13983 \\ 1 & -0.39465 & -0.58060 \\ 1 & 2.03211 & 0 \end{bmatrix} \begin{bmatrix} Y' \\ U \\ V \end{bmatrix}.$$

- Many vision systems only take the Y channel (luminance) as input to reduce computations

Problems





What happened?

- Dimensionality reduction

$$Y = 0.299R + 0.587G + 0.114B$$

- Another view:
 - The conversion is actually a plane equation! All colors on the same plane are converted to the same grayscale value.



Finding a better conversion

- The general form of linear conversion:

$$Y = w_r \cdot R + w_g \cdot G + w_b \cdot B$$

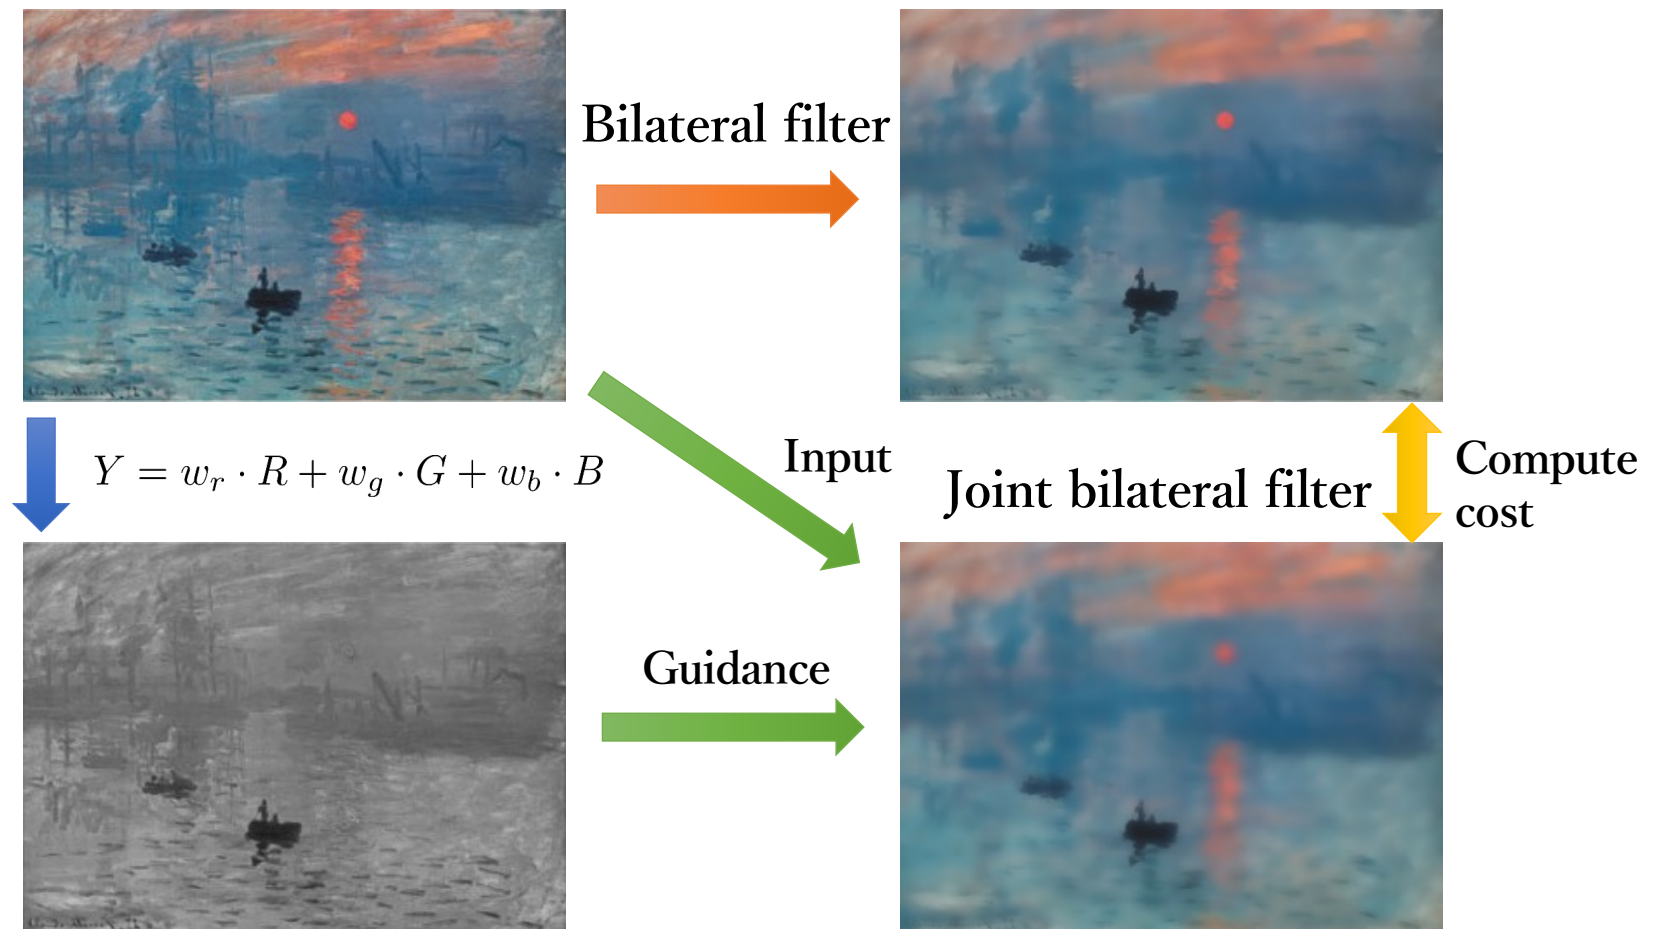
$$w_r, w_g, w_b \geq 0$$

$$w_r + w_g + w_b = 1$$

- Let's consider the quantized weight space $w \in \{0, 0.1, 0.2, \dots, 1\}$
 - For example: $(w_r, w_g, w_b) = (0, 0, 1)$
 $(w_r, w_g, w_b) = (0, 0.1, 0.9)$
 - Given a color image, a set of weight combination corresponds to a grayscale image candidate.
 - We are going to identify which candidate is better!

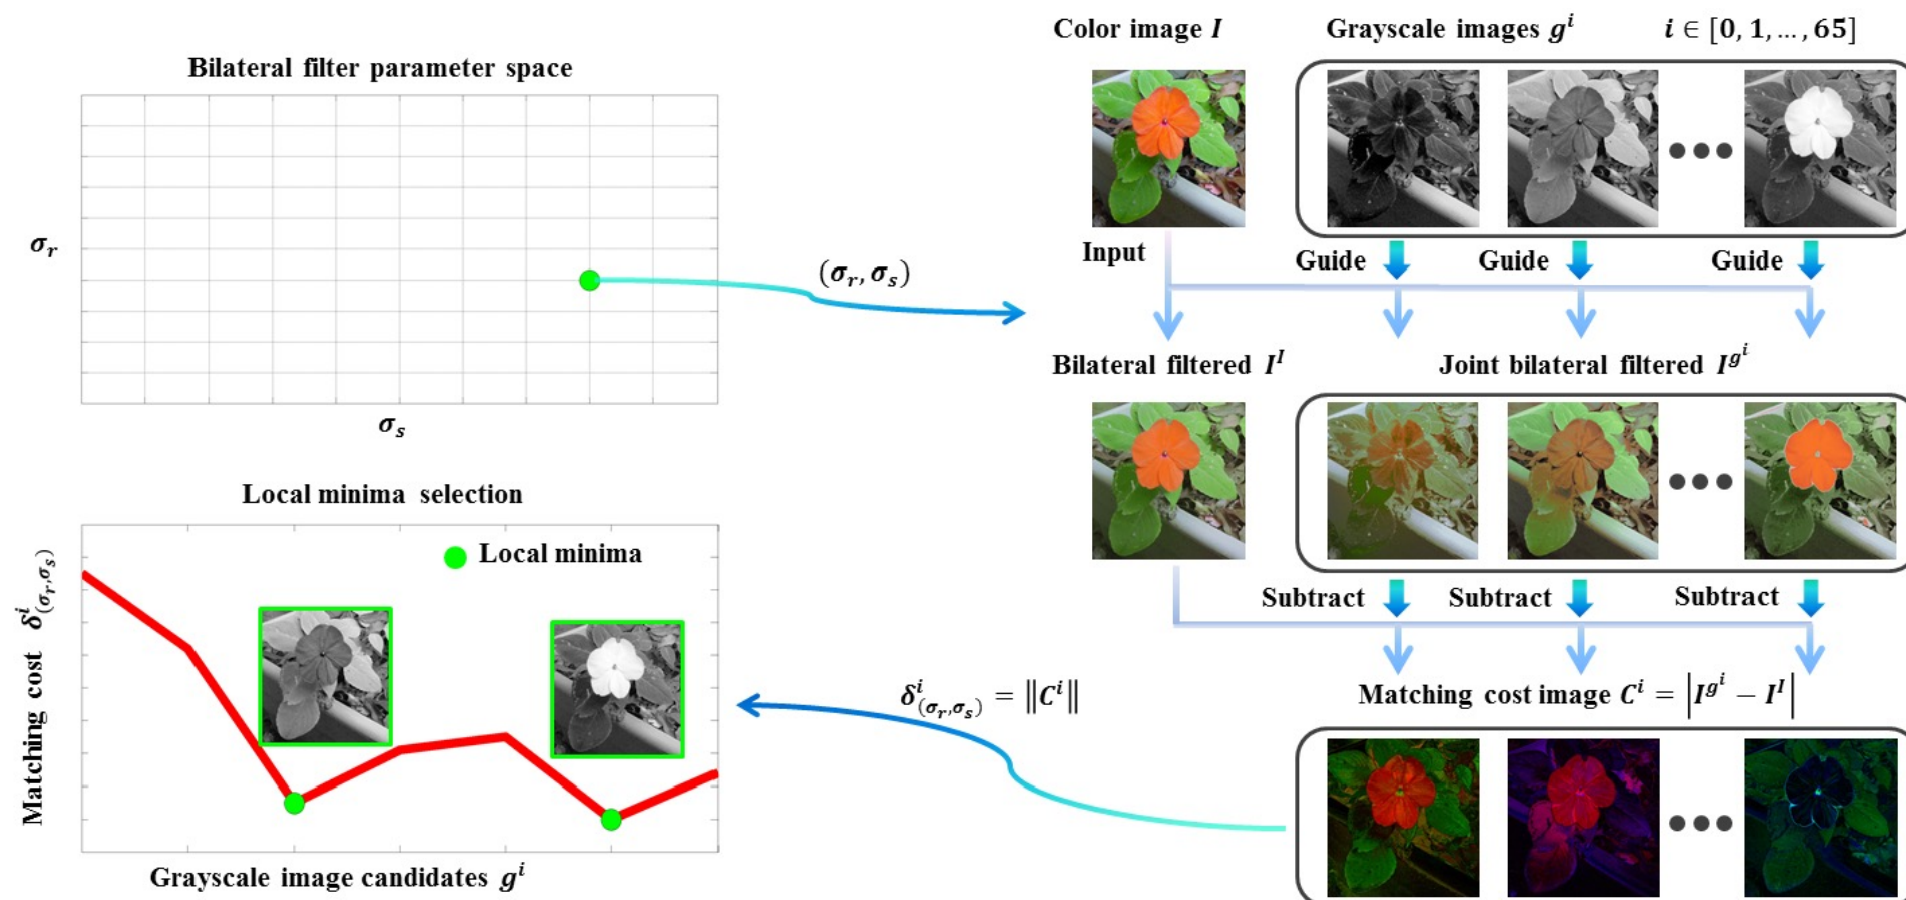
Measuring the perceptual similarity

- Joint bilateral filter (JBF) as the similarity measurement




Measuring the perceptual similarity

- Joint bilateral filter (JBF) as the similarity measurement



Measuring the perceptual similarity

- Find local minimum
 - The actual weight space looks like this:



$$w_r, w_g, w_b \geq 0$$
$$w_r + w_g + w_b = 1$$

Outline

Part 1: Scale Invariant Feature Detection

- Implement Difference of Gaussian (DoG)

Part 2: Image Filtering

- Advanced color-to-gray conversion
- Implement bilateral filter

Rules

- Environment
- Report
- Submission
- Grading
- Late Policy

Bilateral Filter

- Given input image I and guidance T , the bilateral filter is written as:

$$I'_p = \frac{\sum_{q \in \Omega_p} G_s(p, q) \cdot G_r(T_p, T_q) \cdot I_q}{\sum_{q \in \Omega_p} G_s(p, q) \cdot G_r(T_p, T_q)}$$

- I_p : Intensity of pixel p of original image I
- I'_p : Intensity of pixel p of filtered image I'
- T_p : Intensity of pixel p of guidance image T
- Ω_p : Window centered in pixel p
- G_s : Spatial kernel
- G_r : Range kernel

Bilateral Filter

- For the spatial kernel G_s :

$$G_s(p, q) = e^{-\frac{(x_p - x_q)^2 + (y_p - y_q)^2}{2\sigma_s^2}}$$

- For the range kernel G_r :
 - If T is a single-channel image:

$$G_r(T_p, T_q) = e^{-\frac{(T_p - T_q)^2}{2\sigma_r^2}}$$

- If T is a color image:

$$G_r(T_p, T_q) = e^{-\frac{(T_p^r - T_q^r)^2 + (T_p^g - T_q^g)^2 + (T_p^b - T_q^b)^2}{2\sigma_r^2}}$$

- Pixel values should be **normalized to [0, 1]** (divided by 255) to construct range kernel.

Assignment Description

- part2/main.py
 - Read image, execute joint bilateral filter, read setting file, select the best grayscale conversion... etc.
- part2/JBF.py
 - Implement joint bilateral filter

```
class Joint_bilateral_filter(object):
    def __init__(self, sigma_s, sigma_r):
        self.sigma_r = sigma_r
        self.sigma_s = sigma_s
        self.wndw_size = 6*sigma_s+1
        self.pad_w = 3*sigma_s

    def joint_bilateral_filter(self, img, guidance):
        BORDER_TYPE = cv2.BORDER_REFLECT
        padded_img = cv2.copyMakeBorder(img, self.pad_w, self.pad_w, self.pad_w, self.pad_w, BORDER_TYPE)
        padded_guidance = cv2.copyMakeBorder(guidance, self.pad_w, self.pad_w, self.pad_w, self.pad_w, BORDER_TYPE)

        ### TODO ###

        return np.clip(output, 0, 255).astype(np.uint8)
```

Define window size

Pad the input and guidance image

Output image should be in format of uint8

Assignment Description

- part2/eval.py (**DO NOT Modify!**)
 - Evaluate the correctness of the output of joint bilateral filter

```
def main():
    parser = argparse.ArgumentParser(description='evaluation function of joint bilateral filter')
    parser.add_argument('--sigma_s', default=3, type=int, help='sigma of spatial kernel')
    parser.add_argument('--sigma_r', default=0.1, type=float, help='sigma of range kernel')
    parser.add_argument('--image_path', default='./testdata/ex.png', help='path to input image')
    parser.add_argument('--gt_bf_path', default='./testdata/ex_gt_bf.png', help='path to ground truth bf image')
    parser.add_argument('--gt_jbf_path', default='./testdata/ex_gt_jbf.png', help='path to ground truth jbf image')
    args = parser.parse_args()

    img = cv2.imread(args.image_path)
    img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img_gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

    # create JBF class
    JBF = Joint_bilateral_filter(args.sigma_s, args.sigma_r)

    bf_out = JBF.joint_bilateral_filter(img_rgb, img_rgb).astype(np.uint8)
    t0 = time.time()
    jbf_out = JBF.joint_bilateral_filter(img_rgb, img_gray).astype(np.uint8)
    print('[Time] %.4f sec'%(time.time()-t0))
```

We will test your inference duration of joint bilateral filter.

- TAs will run this file to score your code.
- When testing your code, different arguments, e.g. σ_s and σ_r , and corresponding ground truth file will be applied.

Assignment Description

- part2/testdata/
 - One example image with bf and jbf ground truth
 - Two images with respective setting files



1_setting.t

```
R,G,B
0.0,0.0,1.0
0.0,1.0,0.0
0.1,0.0,0.9
0.1,0.4,0.5
0.8,0.2,0.0
sigma_s,2,sigma_r,0.1
```

- There are σ_s , σ_r and 5 groups of gray-conversion parameters in the setting files. Generate **6 gray-scale images** by those parameters and `cv2.cvtColor()`.
- Use those gray-scale images as guidance to run joint bilateral filter and compute the perceptual similarity.
 - Refer p.21 and p.22 for details (use L1-norm as our cost function).
 - Note: casting the image into `np.int32` is needed to avoid overflow when subtraction.

Assignment Description

Recommended steps:

1. Implement joint bilateral filter in JBF.py
2. Use eval.py to evaluate your JBF.py

```
$ python3 eval.py --image_path ./testdata/ex.png --gt_bf_path ./testdata/ex_gt_bf.png  
--gt_jbf_path ./testdata/ex_gt_jbf.png
```

- The error of bilateral and joint bilateral filter should be **both 0**

```
[Error] Bilateral: 0  
[Error] Joint bilateral: 0
```

3. Finish remaining code in main.py if needed
4. Improve the inference speed of joint bilateral filter

Assignment Description

- About the speed test of JBF...
 - Some useful tips
 - Build look-up-table for both spatial and range gaussian kernels
 - Reduce the usage of for-loop to enhance parallel processing
 - Reference time of TA code on ex.png
 - MacBook Pro 2019 (Intel Core i5 CPU) + 8 GB RAM \Rightarrow ~ 1.09 sec
 - Intel Core i7-9700K CPU + 8 GB RAM \Rightarrow ~ 0.93 sec
 - Intel Core i9-9900K CPU + 128 GB RAM \Rightarrow ~ 0.57 sec
 - For fair comparison, you can only use basic functions (e.g. **cannot** use `cv2.filter2D`, `cv2.GaussianBlur`) in JBF.py
 - Cython, multi-thread and GPU acceleration is **forbidden**.

Outline

Part 1: Scale Invariant Feature Detection

- Implement Difference of Gaussian (DoG)

Part 2: Image Filtering

- Advanced color-to-gray conversion
- Implement bilateral filter

Rules

- Environment
- Report
- Submission
- Grading
- Late Policy

Environment

- Package
 - Python == 3.10
 - Numpy == 1.21.6
 - Opencv-python == 4.6.0.66
 - [Python standard library](#)
- You can refer to this [link](#) for environment setup.
- The following APIs are **forbidden** in part2/JBF.py
 - OpenCV: cv2.filter2D, cv2.GaussianBlur
 - Cython, multi-thread and GPU acceleration
- You will **NOT** get corresponding points if you violate rules above.

Report (30%)

- Follow the template we provide and submit pdf file
- Part1: Difference of Gaussian (9%)
 - Plot 8 DoG images described in page.7 with threshold 5.0 on 1.png (4%)
 - Use **three thresholds (2.0, 5.0, 7.0)** on 2.png, plot DoG feature, and describe the difference. (5%)

Report (30%)

- Part2: Joint bilateral filter (21%)
 - For 1.png and 2.png:
 - Report the cost for each filtered image (by using 6 grayscale images as guidance) (1%+1%)
 - Show **original RGB image, two filtered RGB images** and **two grayscale images** with highest and lowest cost (five images in total for each input image) (2%+2%)
 - Describe the difference between those two grayscale images. (5%+5%)
 - Describe how you speed up the implementation of bilateral filter. (5%)

Submission

- Directory architecture:
 - + R12345678/
 - DoG.py
 - JBF.py
- Put above files in a directory (named **StudentID**) and compress the directory into zip file (named **StudentID.zip**)
 - e.g. After TAs run “unzip R12345678.zip”, **it should create one directory named “R12345678”**
- Do NOT copy homework (code and report) from others
- Do NOT show or write images in files you hand in.

Submission

- Please submit those two files, i.e. your **StudentID.zip** and **report.pdf**, separately to **NTU COOL**
- Deadline: **2025/03/27 (Thu.) 23:59**

文件上傳

上傳檔案，或者選擇已上傳的檔案。

 上傳文件

 使用網路攝影機

+ 新增另一個檔案

按一下此處，以找到已上傳的檔案

Grading (Total 15%)

- Part 1 Code: 30%
 - 30%, runs within 5 mins and no error (TA will check your answer on 2.png)
 - 0%, others
- Part 2 Code: 30%
 - 30%, runs within 5 mins and no error (both bf and jbf error == 0)
 - 0%, others
- Report : 30%
- Part 2 Inference time: 10%
 - 10%, Top ~ 20%
 - 6%, 20% ~ 50%
 - 3%, 50% ~ 80%
 - 0%, 80% ~

Late Policy

- Each student has in total 3 days delay quota.
 - Specifically, we take **12 hr. (half day)** as one unit.
 - Each student has **6 units**.
- All four homework assignments share the same 3-day delay.
 - For example, if you delay 1 day in HW1, then you will have 2 day delay quota left for HW2, HW3 and HW4.
- If you have spent all 6 delay units, your late submission will not be accepted.
 - That is, you are required to submit your homework in time to get points.
- If you have questions about the policy, you can contact TAs.

TA information

- Zih-Syuan Lin(林子軒)
E-mail: zslin@media.ee.ntu.edu.tw
TA hour: Mon. 13:30 - 15:30
Location: 博理 421
- Yu-Ching Fan (范宇清)
E-mail: jackmafan@media.ee.ntu.edu.tw
Location: 博理 421